

Salary Prediction Classification

A data science project by Rafif





About this project

In this project, I make machine learning models to predict whether a person makes $> \$50K$ or $\leq \$50K$ a year

The dataset was obtained from [Kaggle](#). It was acquired from the 1994 Census database by Barry Becker.

This project focuses on: data pre-processing, EDA, data visualization, feature selection, and classification algorithm.



Dataset

This dataset contains 15 columns:

- 6 numerical columns
- 8 categorical columns
- 1 target column (salary)

Data columns (total 15 columns):				
#	Column	Non-Null Count		Dtype
0	age	32561	non-null	int64
1	workclass	32561	non-null	object
2	fnlwgt	32561	non-null	int64
3	education	32561	non-null	object
4	education-num	32561	non-null	int64
5	marital-status	32561	non-null	object
6	occupation	32561	non-null	object
7	relationship	32561	non-null	object
8	race	32561	non-null	object
9	sex	32561	non-null	object
10	capital-gain	32561	non-null	int64
11	capital-loss	32561	non-null	int64
12	hours-per-week	32561	non-null	int64
13	native-country	32561	non-null	object
14	salary	32561	non-null	object

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

Pre-processing

Checking for any missing value (NaN)

```
1 # Check the number of NaN
2 # values in every column
3 df.isna().sum()
```

age	0
workclass	0
fnlwgt	0
education	0
education-num	0
marital-status	0
occupation	0
relationship	0
race	0
sex	0
capital-gain	0
capital-loss	0
hours-per-week	0
native-country	0
salary	0

No missing values? Sometimes they are marked as string characters (e.g. empty space ' ')!

Column: workclass

Private	22696
Self-emp-not-inc	2541
Local-gov	2093
?	1836
State-gov	1298
Self-emp-inc	1116
Federal-gov	960
Without-pay	14
Never-worked	7

They are marked as '?'



Pre-processing

Replacing all '?'s with NaNs → there are 2399 (~7.4%) rows with missing values

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
14	40	Private	121772	Assoc-voc	11	Married-civ-spouse	Craft-repair	Husband	Asian-Pac-Islander	Male	0	0	40	NaN	>50K
27	54	NaN	180211	Some-college	10	Married-civ-spouse	NaN	Husband	Asian-Pac-Islander	Male	0	0	60	South	>50K
38	31	Private	84154	Some-college	10	Married-civ-spouse	Sales	Husband	White	Male	0	0	38	NaN	>50K
51	18	Private	226956	HS-grad	9	Never-married	Other-service	Own-child	White	Female	0	0	30	NaN	<=50K
61	32	NaN	293936	7th-8th	4	Married-spouse-absent	NaN	Not-in-family	White	Male	0	0	40	NaN	<=50K

These rows are dropped altogether



Pre-processing

45 rows have duplicates

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
2303	90	Private	52386	Some-college	10	Never-married	Other-service	Not-in-family	Asian-Pac-Islander	Male	0	0	35	United-States	<=50K
3917	19	Private	251579	Some-college	10	Never-married	Other-service	Own-child	White	Male	0	0	14	United-States	<=50K
4325	25	Private	308144	Bachelors	13	Never-married	Craft-repair	Not-in-family	White	Male	0	0	40	Mexico	<=50K
4767	21	Private	250051	Some-college	10	Never-married	Prof-specialty	Own-child	White	Female	0	0	10	United-States	<=50K
4881	25	Private	308144	Bachelors	13	Never-married	Craft-repair	Not-in-family	White	Male	0	0	40	Mexico	<=50K

Dropped the duplicates and kept the first row



Pre-processing

Binary encoding for salary

```
1 df['salary'] = df['salary'].apply(lambda x: 0 if x == ' <=50K' else 1)
2
3 df.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	salary
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	0
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	0
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	0
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	0
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	0



Train-test split

To test the performance of my model-to-be with unseen data, I split the data into training and test sets with a ratio of 80:20.

```
from sklearn.model_selection import train_test_split


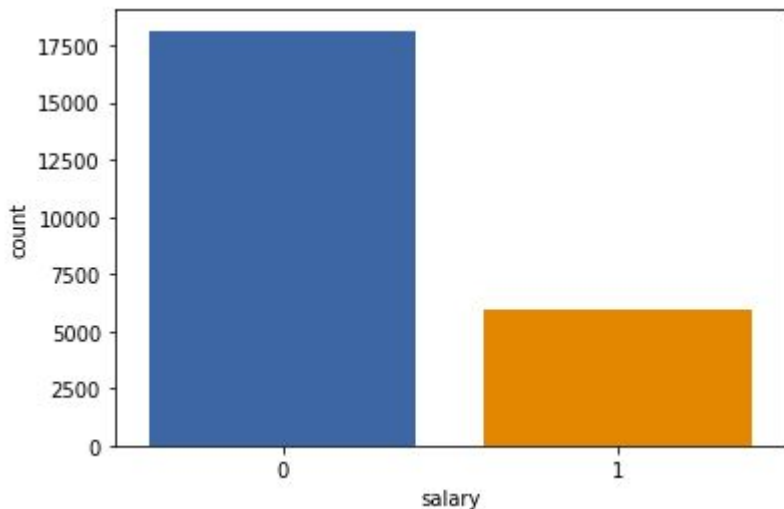
# Select the features
X = df.drop('salary',axis=1)
# Select the target
y = df[['salary']]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```




Data Understanding

The target feature is not balanced (`salary=1` is ~3 times `salary=0`)

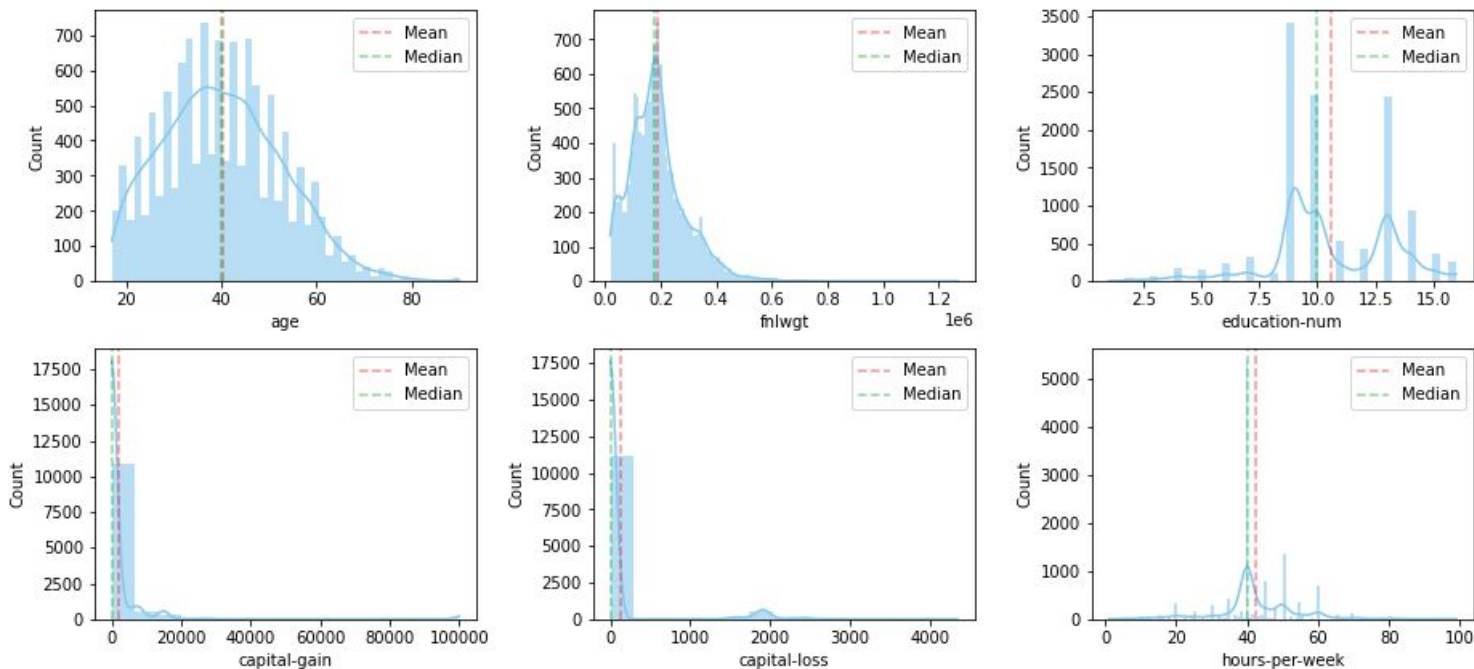


Undersampling data with `salary = 0`



Data Understanding

Distribution of numerical features:



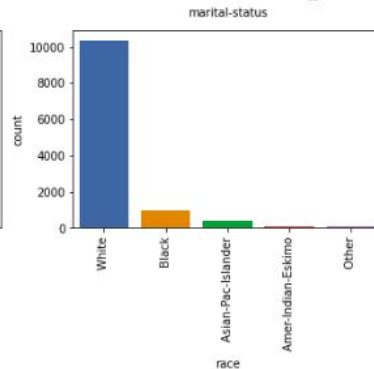







A lot of zeros in
capital-gain
and
capital-loss

I kept them as
they are



Data Understanding


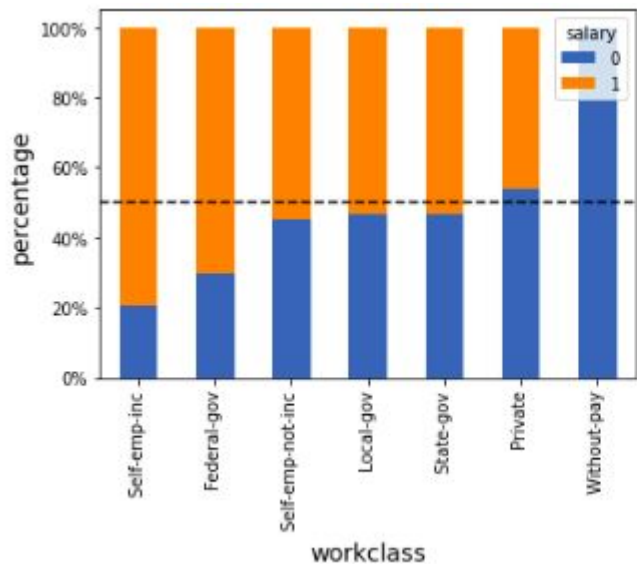
Count plot of categorical features:





Data Understanding

The percentage of people based on their salary in each categorical value for `workclass` and `education`:



Self-employed people and doctoral graduates tend to make more salary

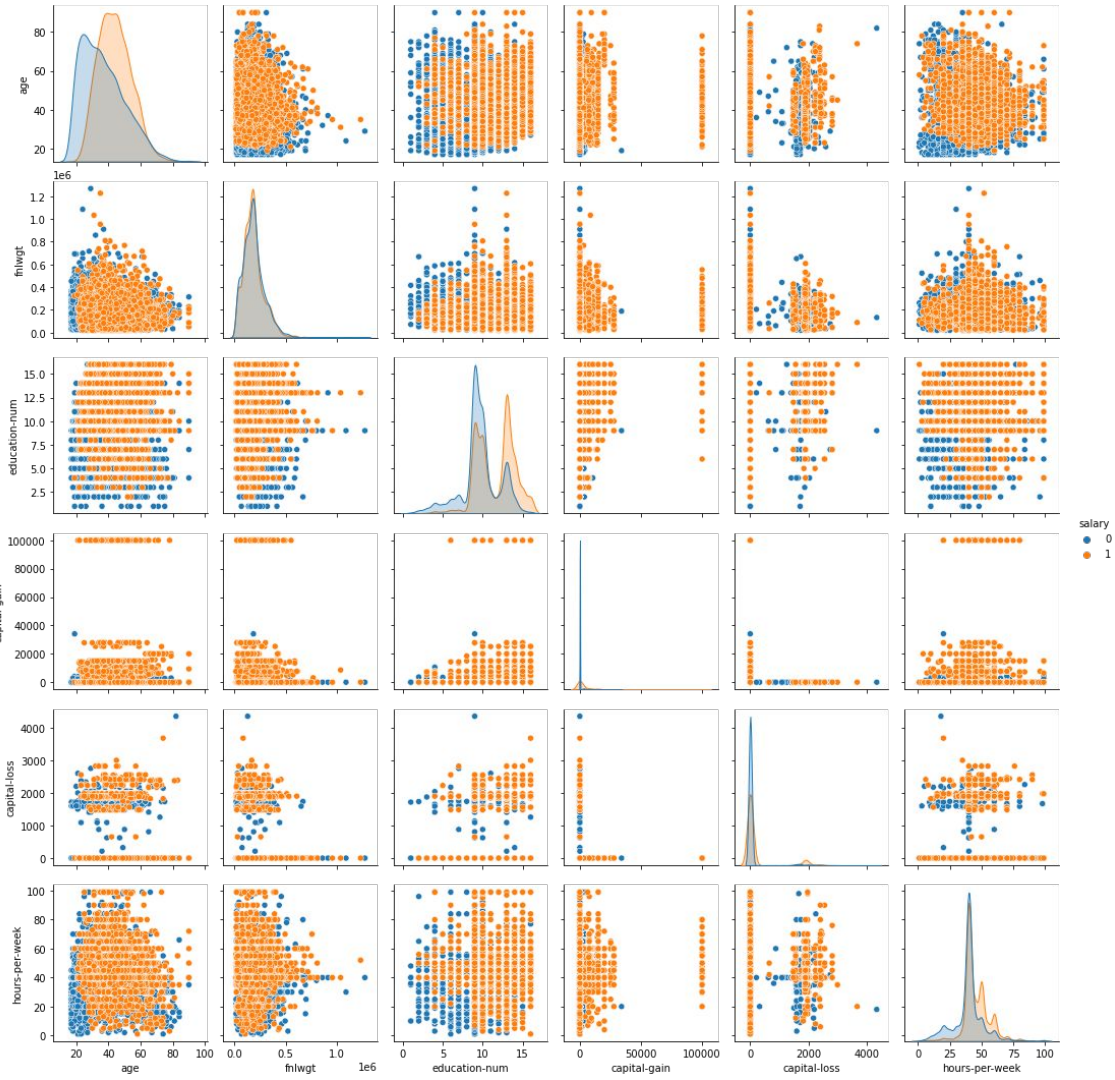
Overall there are correlations with salary



Data Understanding

No multicollinearity among numerical features

No clear distinction in all features based on salary.





Building the Models

- Not dropping any feature (the numerical features are weakly correlated with `salary` anyway)
- Classification algorithms: logistic regression and random forest
- Grid search with cross-validation to tune the parameters

Building the Models: logistic regression

Prerequisite I: label encoding using sklearn

```
# We first apply one hot encoding for all categorical features
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder

transformer = ColumnTransformer(transformers=[('cat', OneHotEncoder(drop='first',
                                                                    handle_unknown='ignore'),
                                                                    cat_col)],
                                ,remainder='passthrough')

transformer.fit(X_train)
# Label encoding for training data
X_train_encoded = pd.DataFrame(transformer.transform(X_train).toarray(), columns=transformer.get_feature_names_out())
# Label encoding for test data
X_test_encoded = pd.DataFrame(transformer.transform(X_test).toarray(), columns=transformer.get_feature_names_out())
```

Prerequisite II: scale the features (to interpret the coefficients easier)

```
# To check feature importance later, we use the same scale across all features
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
scaler.fit(X_train_encoded)

X_train_scaled = scaler.transform(X_train_encoded)
X_test_scaled = scaler.transform(X_test_encoded)
```


Building the Models: logistic regression

```
# Import the logistic regression model from sklearn
from sklearn.linear_model import LogisticRegression

# Import the grid search algorithm from sklearn
from sklearn.model_selection import GridSearchCV

import warnings
warnings.filterwarnings('ignore')

# parameter grid to be tested for the logistic regression model
parameters = {
    'penalty' : ['l1', 'l2'],
    'C'       : np.logspace(-2, 2, 5),
    'solver'  : ['newton-cg', 'lbfgs', 'liblinear'],
}

logreg = LogisticRegression()
clf = GridSearchCV(logreg,
                   param_grid=parameters,
                   scoring='accuracy',
                   cv=5)

clf.fit(X_train_scaled, y_train)
```

Grid search method to find the best parameters

Hyperparameters:

- Penalty
- C
- solver

```
# Print the best model and its accuracy from the cross-validation method
print("Tuned Hyperparameters :", clf.best_params_)
print("Accuracy :", clf.best_score_)
```

Tuned Hyperparameters : {'C': 1.0, 'penalty': 'l1', 'solver': 'liblinear'}
Accuracy : 0.821231150966241



Building the Models: logistic regression

	feature	feature_importance
15	cat__education_ Doctorate	4.021244
19	cat__education_ Prof-school	4.015817
17	cat__education_ Masters	3.115950
21	cat__marital-status_ Married-AF-spouse	3.066734
14	cat__education_ Bachelors	2.618511
22	cat__marital-status_ Married-civ-spouse	2.480619
12	cat__education_ Assoc-acdm	2.004323
13	cat__education_ Assoc-voc	1.845402
44	cat__relationship_ Wife	1.585974
20	cat__education_ Some-college	1.397390

Consistent with the results from EDA!

Building the Models: logistic regression

```
print(confusion_matrix(y_test,y_pred),'\n')
print(classification_report(y_test,y_pred))
```

```
[[4308  177]
 [ 934  609]]
```

	precision	recall	f1-score	support
0	0.82	0.96	0.89	4485
1	0.77	0.39	0.52	1543
accuracy			0.82	6028
macro avg	0.80	0.68	0.70	6028
weighted avg	0.81	0.82	0.79	6028

```
print(f'accuracy = {accuracy_score(y_test,y_pred)}')
print(f'ROC AUC score = {roc_auc_score(y_test,y_pred)}')
```

```
accuracy = 0.8156934306569343
ROC AUC score = 0.6776103971544813
```

The accuracy is 81.6%, but the AUC score is only 68.8%

Since the test salary is also very likely to be imbalanced, this indicates that this model is not good enough

Building the Models: random forest

```
# Import random forest from sklearn
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state=7)
parameters = {'n_estimators':[100,200,300,400,500],
              'criterion': ['gini','entropy']}

clf = GridSearchCV(estimator=rf,
                   param_grid=parameters,
                   scoring='accuracy',
                   cv=5)

clf.fit(X_train_encoded,y_train)
```

Grid search method to find the best parameters

Hyperparameters:

- n_estimators
- criterion

```
# Print the best model and its accuracy from the cross-validation method
print("Tuned Hyperparameters :", clf.best_params_)
print("Accuracy :",clf.best_score_)
```

```
Tuned Hyperparameters : {'criterion': 'entropy', 'n_estimators': 300}
Accuracy : 0.8181282147256621
```

Building the Models: random forest

```
print(confusion_matrix(y_test,y_predrf),'\n')
print(classification_report(y_predrf, y_test))
```

```
[[3632  853]
 [ 241 1302]]
```

	precision	recall	f1-score	support
0	0.81	0.94	0.87	3873
1	0.84	0.60	0.70	2155
accuracy			0.82	6028
macro avg	0.83	0.77	0.79	6028
weighted avg	0.82	0.82	0.81	6028

```
print(f'accuracy = {accuracy_score(y_test,y_predrf)}')
print(f'ROC AUC score = {roc_auc_score(y_test,y_predrf)}')
```

```
accuracy = 0.818513603185136
ROC AUC score = 0.8268106188194103
```

The accuracy and AUC score are both >80%, indicating this model performs quite well.

This is expected since the random forest algorithm is more complex than logistic regression



Conclusions

1. This dataset contains 15 columns, 6 are numerical, 8 are categorical, and the target `salary` is binary.
2. The numerical features are very weakly correlated with `salary`, and the categorical features seem to have stronger correlations with `salary`.
3. I use two classification algorithms: logistic regression and random forest to predict `salary`.
4. The best model is the random forest classifier, where I can obtain an accuracy of 81.8%, and an ROC AUC score of 82.7%.